

Data Preprocessing



Hello!

I am Eslam Ahmed

I am a software engineer.

You can find me at jeksogsa@gmail.com



Hello!

I am Eman Ehab

I am a ML research engineer.

You can find me at
emanehab.ieee@gmail.com



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Agenda

What is Data

Machine Learning

Data Preprocessing

Feature Engineering and Extraction

- Domain knowledge features
- Date and Time features
- String operations
- Web Data
- Geospatial features
- Work with Text
- Work with Images
- Work with Audio
- Dimension Reduction with PCA

Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

Feature Selection

What is Data

We are living in a data-driven economy. It's a world where having tons of data, understanding it and knowing what to do with data is power.

Understanding your data is not one of the most difficult things in data science, but it is time-consuming.
Interpretation of data is effective when we know about the source of data.



What is Data (Types of Data)

Categorical

This represents qualitative data with no apparent inherent mathematical meaning.

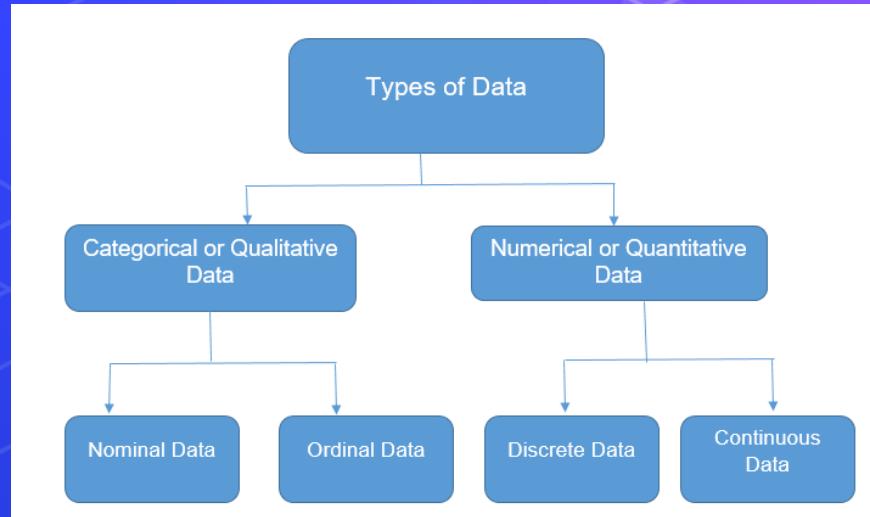
Example: Yes or No, Sex, Race, Marital Status etc.

These can be assigned numbers like Yes(0) and No(1), but numbers have no mathematical meaning.

Numerical

This represents some sort of quantitative measurement.

Example: height of people, stock price, page load time etc



What is Data (Numerical)

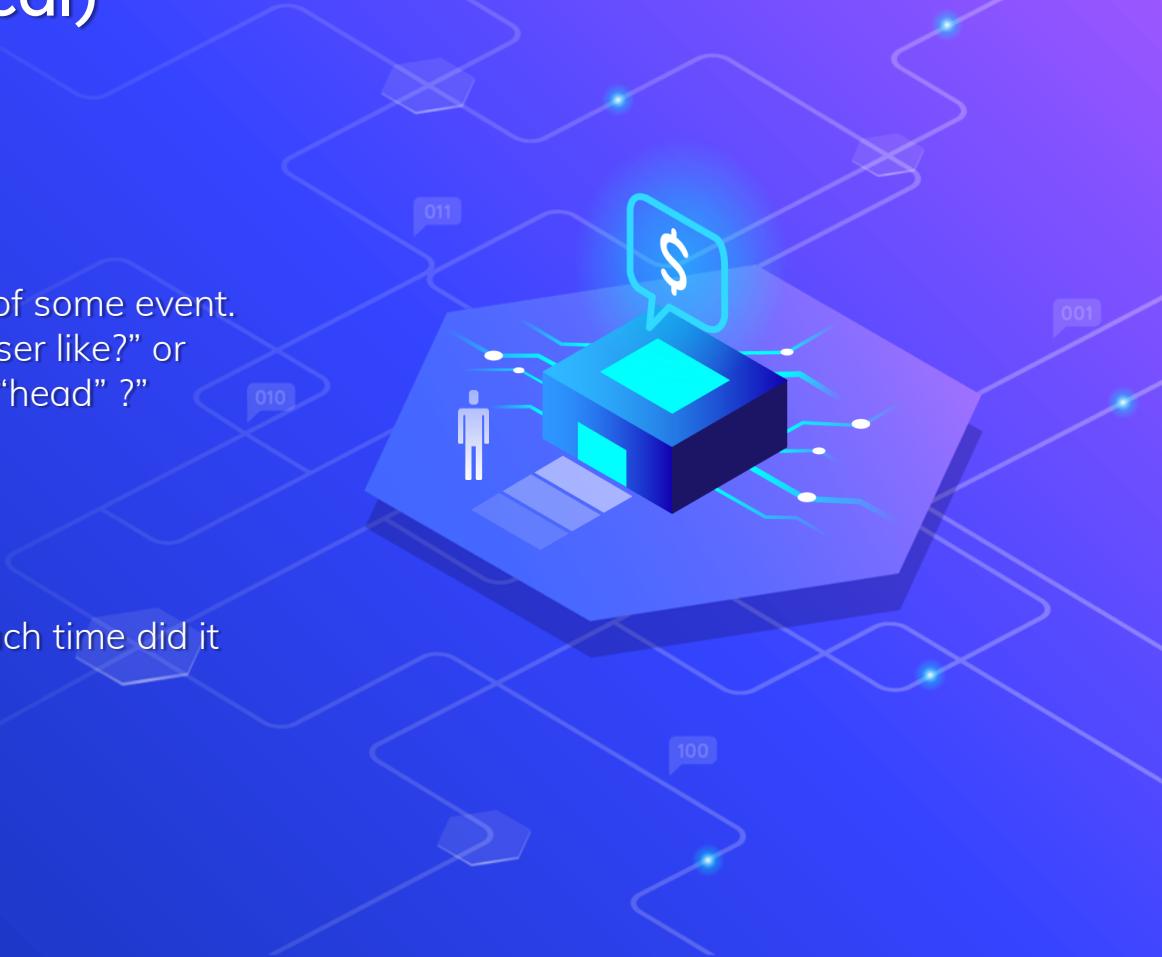
Discrete data

This is integer based, often counts of some event.

Example: "How many songs do a user like?" or
"How many times a coin flipped to "head" ?"

Continuous data

It has an infinite number of
possible values. Example: "How much time did it
take for a user to buy?"



What is Data (Categorical)

Nominal or unordered data

we assign individual items to named categories that do not have an implicit or natural value or rank.

Example: Animals can be cats or dogs or horses etc.. that would be nominal data.

Ordered or ordinal data

items are assigned to categories that do have some kind of implicit or natural order.

Example: "Small, Medium, or Large." Another example is a survey question that asks us to rate an item on a 1 to 10 scale, with 10 being the best. This implies that 10 is better than 9, which is better than 8, and so on.



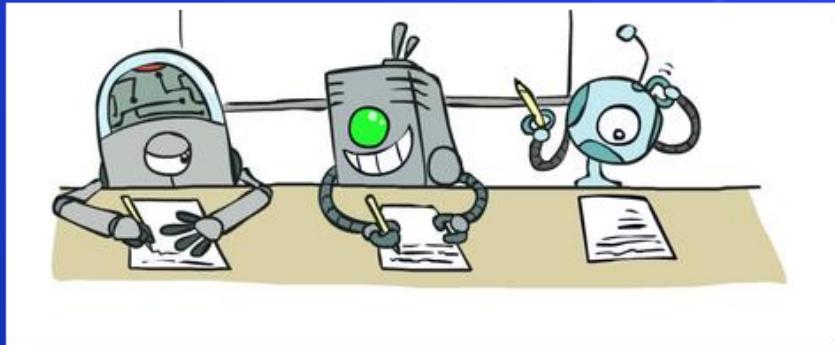
Agenda

- hexagon icon What is Data
- hexagon icon **Machine Learning**
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

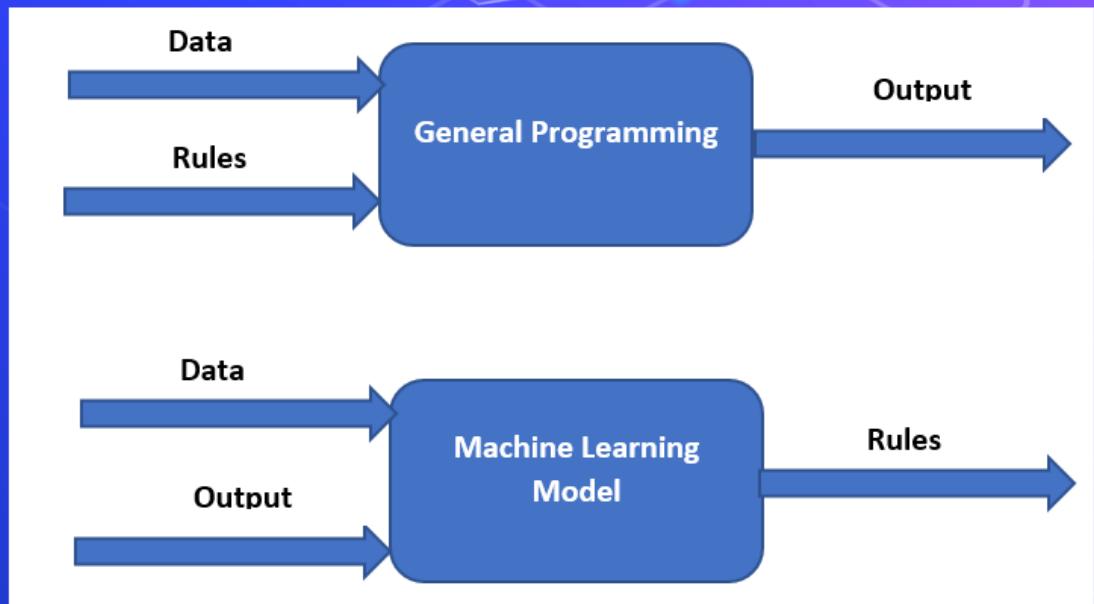
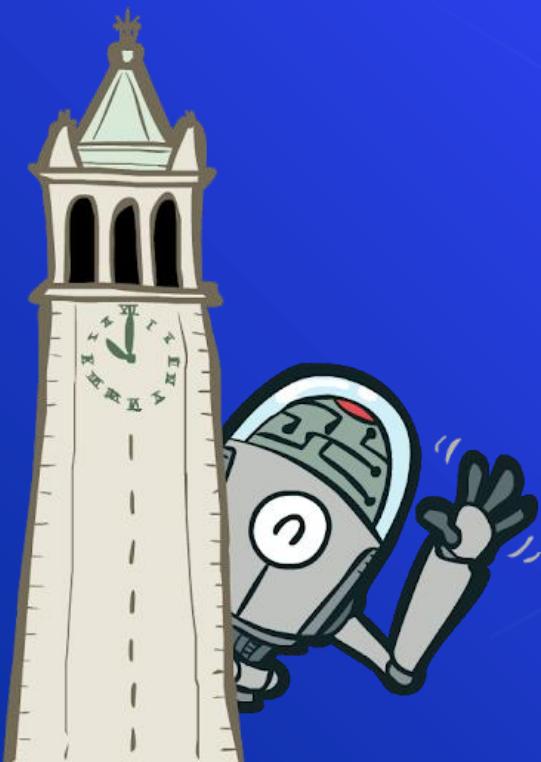
- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Machine Learning

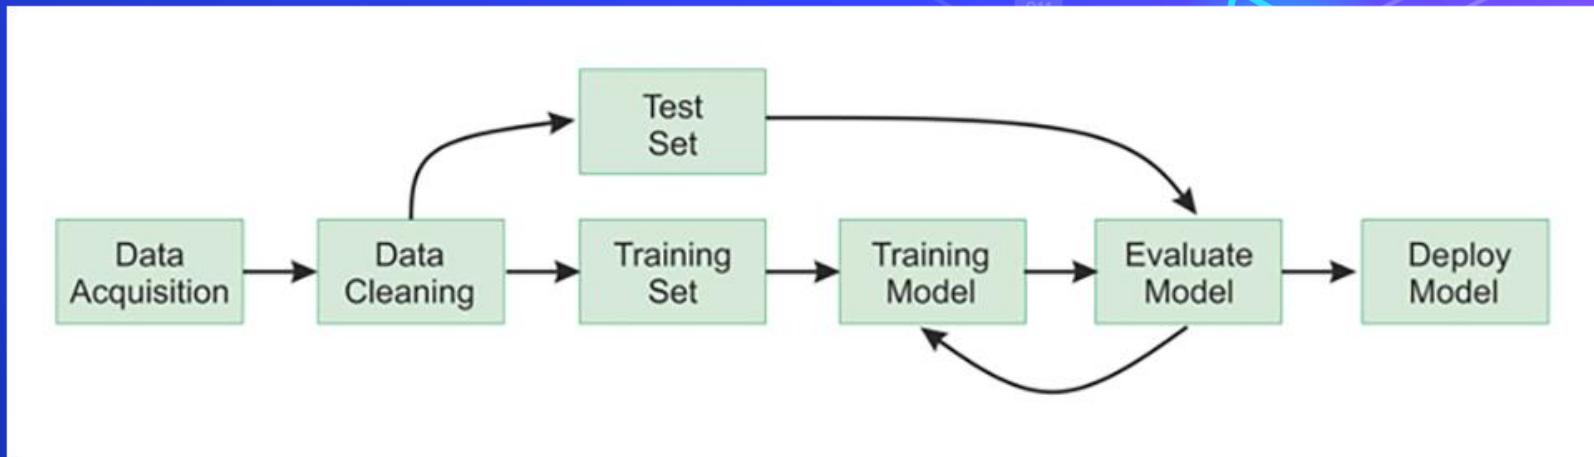
Machine learning involves computers discovering how they can perform tasks without being explicitly programmed to do so. It involves computers learning from data provided so that they carry out certain tasks.



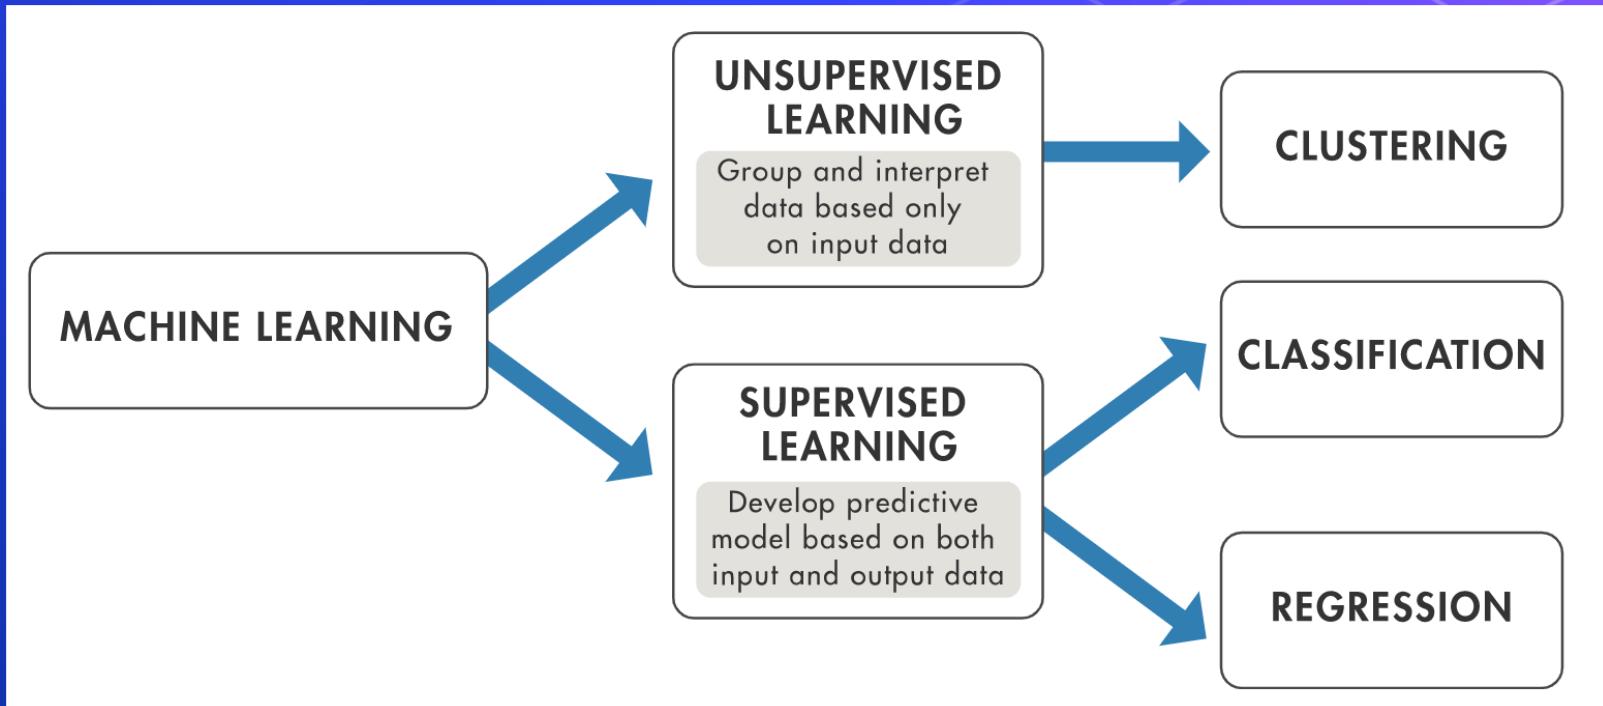
Machine Learning



Machine Learning



Machine Learning



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon **Data Preprocessing**
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

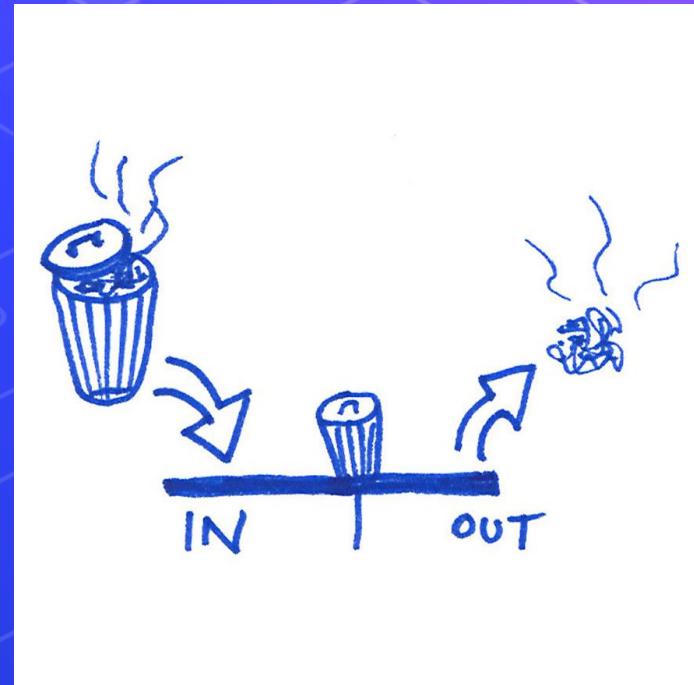
- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Data Preprocessing

The phrase "garbage in, garbage out" is the best description for the use of data preprocessing.

Data gathering methods are often loosely controlled, resulting in out-of-range values, useless features, outliers, missing values, etc.

Garbage data produce misleading results. Thus, data preprocessing is the most important phase of a machine learning project.



Data Preprocessing

Feature Engineering and Extraction

Extract and Create useful features from raw data into features suitable for modeling.

Feature Transformation

Transformation of data to improve the accuracy of the algorithm.

Feature Selection

Removing unnecessary features.



Data Preprocessing

Install

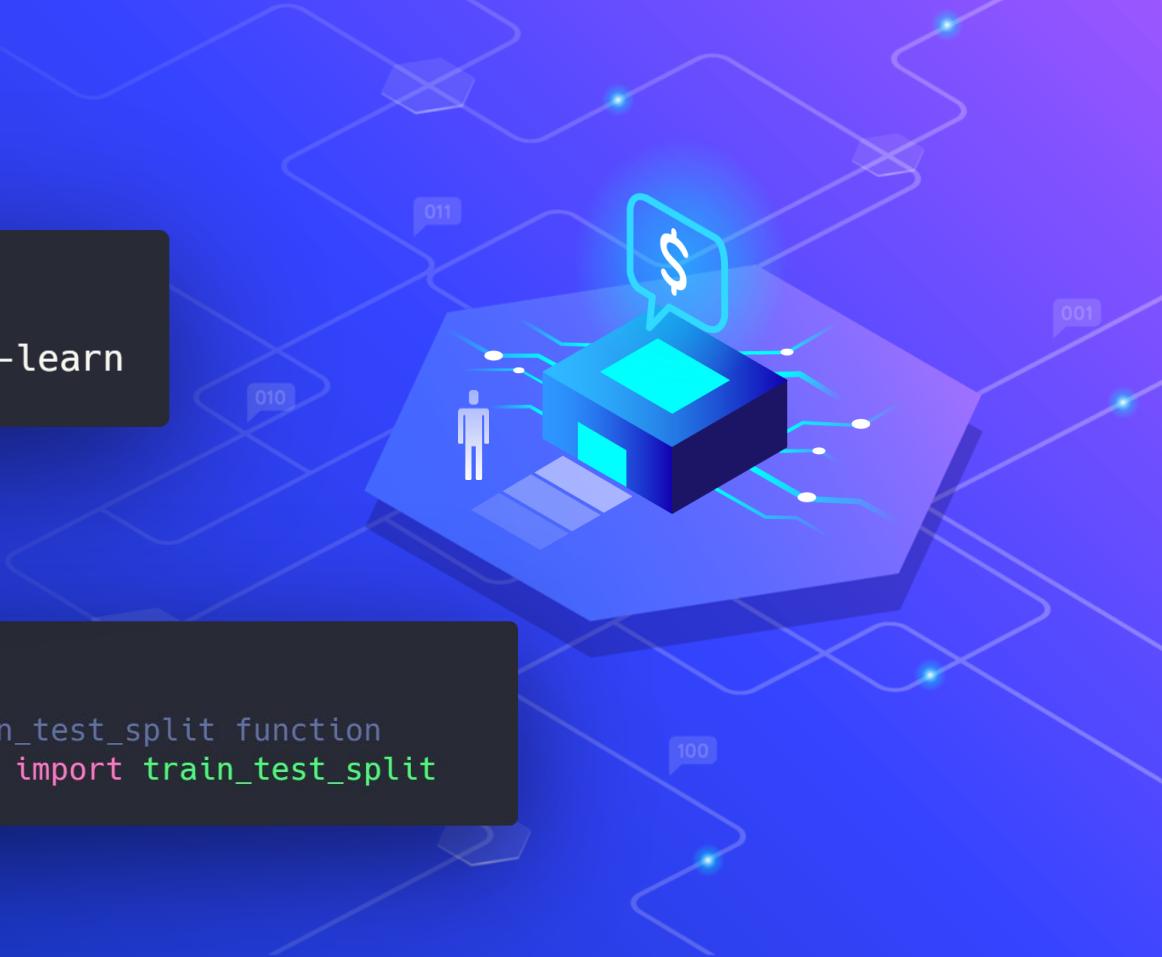


```
1 >_ conda install scikit-learn
```

Use



```
1 # example for importing train_test_split function  
2 from sklearn.model_selection import train_test_split
```



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Domain knowledge features

It's the process of creating new useful features from current features based on our understanding of the problem domain knowledge.



```
1 # example for delivery system we calculate speed in Km/m  
2 df['speed'] = df['Distance (KM)'] / (df['Time from Pickup to Arrival'] / 60)
```

Agenda



What is Data



Machine Learning



Data Preprocessing



Feature Engineering and Extraction

- Domain knowledge features
- **Date and Time features**
- String operations
- Web Data
- Geospatial features
- Work with Text
- Work with Images
- Work with Audio
- Dimension Reduction with PCA



Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines



Feature Selection

Date and Time features

Date columns usually provide valuable information about the problem, they are neglected as an input for the machine learning algorithms. It might be the reason for this, that dates can be present in numerous formats, which make it hard to understand by algorithms.



```
1 df['Time'] = pd.to_datetime(df['Time'],format='%d/%m/%Y %H:%M')
```

Date and Time features

```
1 df['Year'] = df['Time'].dt.year  
2 df['Month'] = df['Time'].dt.month  
3 df['Month_Name'] = df['Time'].dt.month_name()  
4 df['Week'] = df['Time'].dt.week  
5 df['Day'] = df['Time'].dt.day  
6 df['Week_Day'] = df['Time'].dt.weekday  
7 df['Day_Name'] = df['Time'].dt.day_name()  
8 df['Hour'] = df['Time'].dt.hour  
9 df['Minute'] = df['Time'].dt.minute
```

- Extract day, hour, minute, seconds, quarter, month, year, etc.
- Extract time-based features like evenings, noon, night time etc.
- Extract seasonal features like winter, summer, autumn.
- Calculate time elapsed between two related Date features.



Date and Time features

```
1 def map_hours(x):
2     if x in [0,1,2,3,4,5,6,7,8,9,10,11,12]:
3         return 'morning'
4     elif x in [13,14,15,16]:
5         return 'afternoon'
6     else:
7         return 'evening'
8
9 df['Period'] = df['Hour'].apply(map_hours)
```

- Extract day, hour, minute, seconds, quarter, month, year, etc.
- **Extract time-based features like evenings, noon, night time etc.**
- Extract seasonal features like winter, summer, autumn.
- Calculate time elapsed between two related Date features.



Date and Time features

```
1 def map_months(x):
2     if x in [12, 1, 2]:
3         return 'Winter'
4     elif x in [3, 4, 5]:
5         return 'Spring'
6     elif x in [6, 7, 8]:
7         return 'Summer'
8     elif x in [9, 10, 11]:
9         return 'Autumn'
10
11 df['Season'] = df['Month'].apply(map_months)
```

- Extract day, hour, minute, seconds, quarter, month, year, etc.
- Extract time-based features like evenings, noon, night time etc.
- **Extract seasonal features like winter, summer, autumn.**
- Calculate time elapsed between two related Date features.

Date and Time features



```
1 df['Elapsed_Years'] = (datetime.now() - df['Time']) / np.timedelta64(1, 'Y')
```

- Extract day, hour, minute, seconds, quarter, month, year, etc.
- Extract time-based features like evenings, noon, night time etc.
- Extract seasonal features like winter, summer, autumn.
- Calculate time elapsed between two related Date features.

Agenda



What is Data

Machine Learning

Data Preprocessing

Feature Engineering and Extraction

- Domain knowledge features
- Date and Time features
- String operations
- Web Data
- Geospatial features
- Work with Text
- Work with Images
- Work with Audio
- Dimension Reduction with PCA



Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines



Feature Selection

String operations

There is a lot of data hidden within strings, you can use a lot of string operations techniques to extract those data.

```
1 def extract_email_provider(email):
2     return email.split('@')[1]
3
4 df['Email_Provider'] = df['Email'].apply(extract_email_provider)
```

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Web Data

There are a lot of web applications data logs like user agents that contain a lot of info about used browser, OS and device.

```
1 import user_agents
2
3 ua = 'Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
      Chrome/83.0.4103.116 Safari/537.36'
4 ua = user_agents.parse(ua)
5
6 print('Is a bot? ', ua.is_bot)
7 print('Is mobile? ', ua.is_mobile)
8 print('Is PC? ', ua.is_pc)
9 print('OS Family: ', ua.os.family)
10 print('OS Version: ', ua.os.version)
11 print('Browser Family: ', ua.browser.family)
12 print('Browser Version: ', ua.browser.version)
13 print('Device Family: ', ua.device.family)
14 print('Device Brand: ', ua.device.brand)
15 print('Device Model: ', ua.device.model)
```

Web Data

IP address is a great way to extract geographical data from it.

```
1 from ip2geotools.databases.noncommercial import DbIpCity as ip2geo
2
3
4 response = ip2geo.get('45.243.72.231', api_key='free')
5
6 print(response.ip_address)
7 print(response.city)
8 print(response.region)
9 print(response.country)
10 print(response.latitude)
11 print(response.longitude)
```

Agenda



What is Data



Machine Learning



Data Preprocessing



Feature Engineering and Extraction

- Domain knowledge features
- Date and Time features
- String operations
- Web Data
- Geospatial features
- Work with Text
- Work with Images
- Work with Audio
- Dimension Reduction with PCA



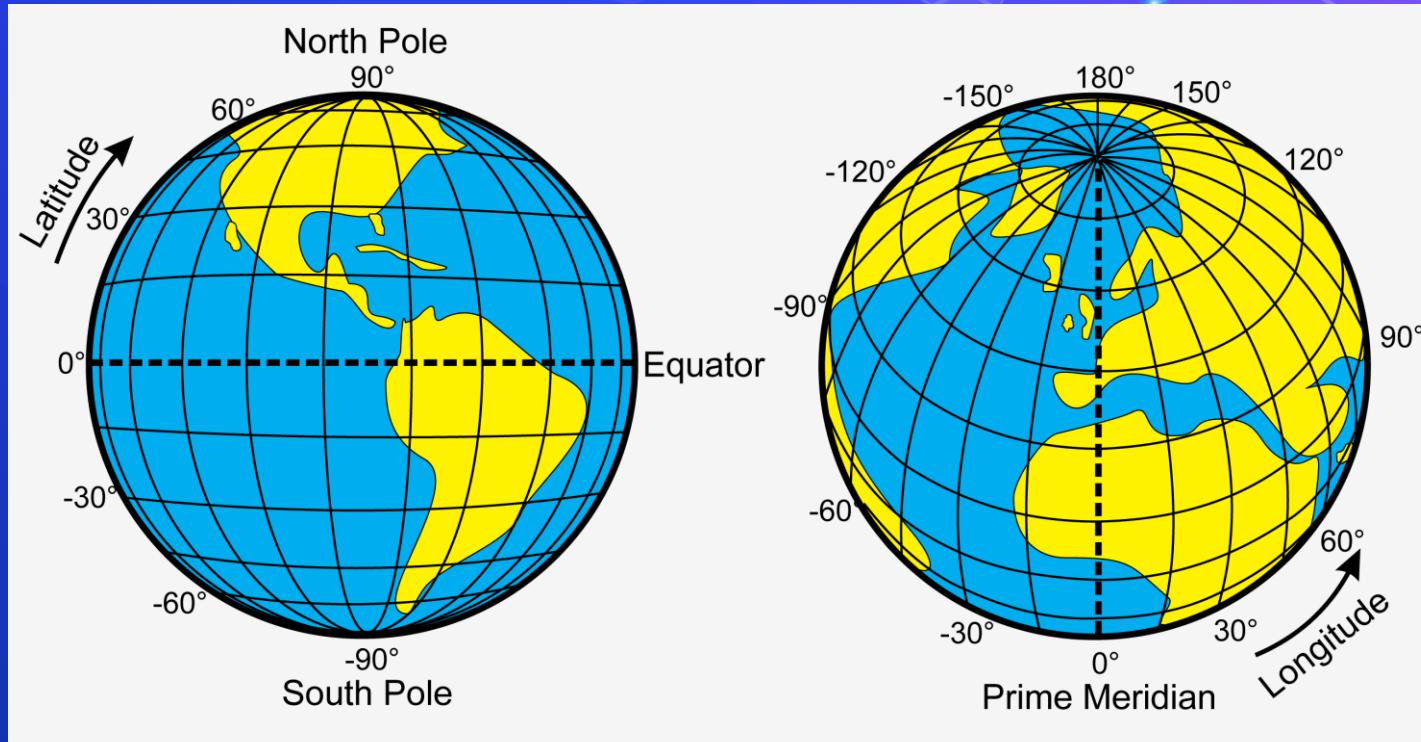
Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines



Feature Selection

Geospatial features

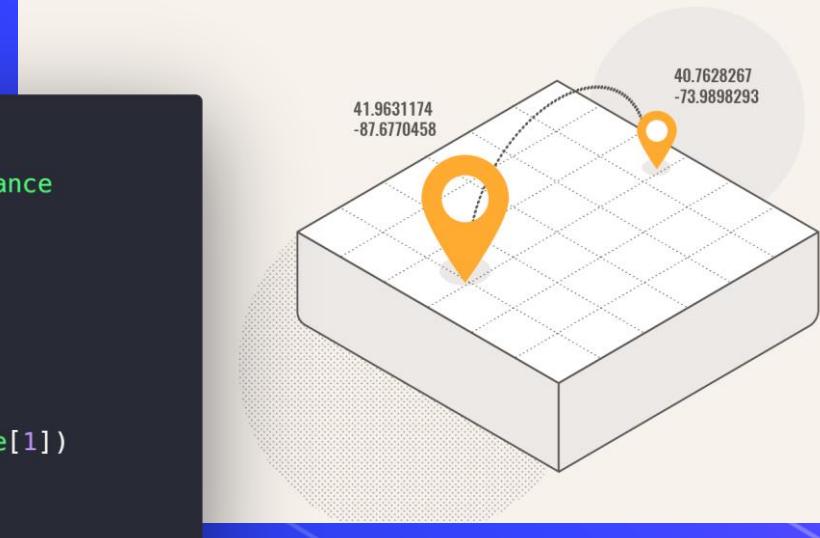


Geospatial features

Measure distance with `haversine_distance`

<https://www.movable-type.co.uk/scripts/latlong.html>

```
1 from datasist.feature_engineering import haversine_distance
2
3
4 # lat, long
5 my_home = [30.109919, 31.308797]
6 cafe = [30.120982, 31.322026]
7
8
9 haversine_distance(my_home[0], my_home[1], cafe[0], cafe[1])
10 """
11 1.769848
12 """
```



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - **Work with Text**
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Work with Text

BAG OF WORDS

Converts text to a matrix where every row is an observation and every feature is a unique word. The value of each element in the matrix is either a binary indicator marking the presence of that word or an integer of the number of times that word appears.

Chris Albon



```
1 from sklearn.feature_extraction.text import CountVectorizer  
2  
3 vectorizer = CountVectorizer(stop_words='english')  
4 txt_feats = vectorizer.fit_transform(text)  
5 txt_feats.to_array()
```

Work with Text

TF-IDF

TF-IDF is a measure of originality of a word by comparing the number of times a word appears in a doc with the number of docs the word appears in.

$$\text{TF-IDF} = \text{TF}(t, d) \times \text{IDF}(t)$$

Term frequency
Number of times term t appears in a doc, d

$$\log \frac{1 + n}{1 + df(d, t)}$$

Document frequency of the term t

Inverse document frequency
 $\#$ documents

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2
3 vectorizer = TfidfVectorizer(stop_words='english')
4 txt_feats = vectorizer.fit_transform(text)
5 txt_feats.toarray()
```

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - **Work with Images**
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Work with Images

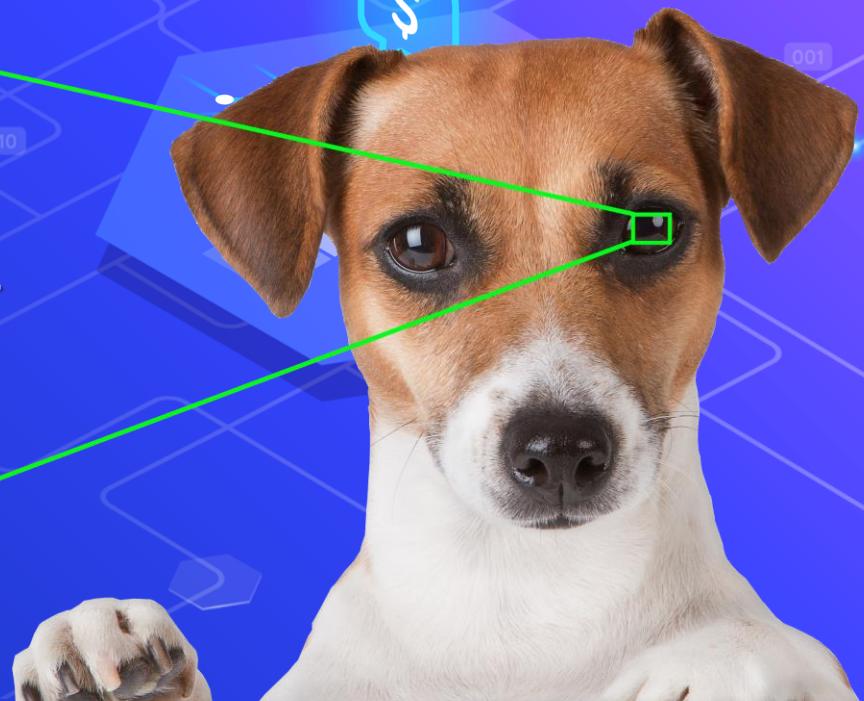
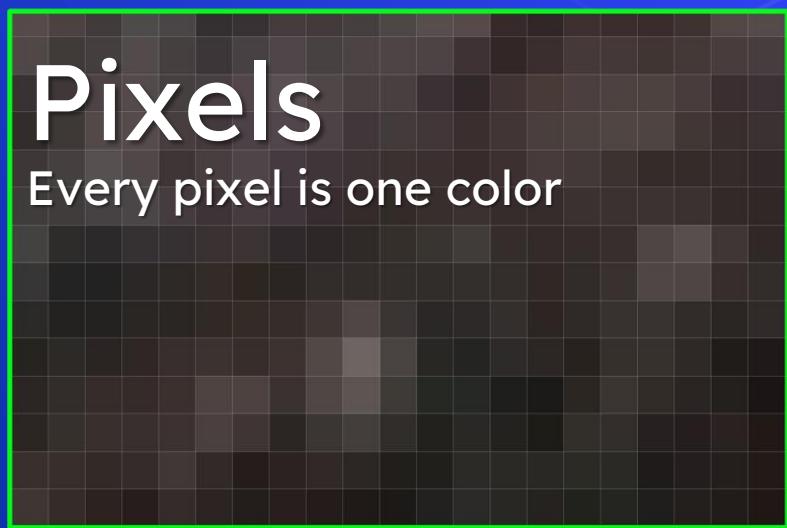


What We See

08 02 22 97 38 15 00 40 00 75 04 05 07 78 52 12 50 77 91 08 08 02 22 97
49 49 99 40 17 81 18 57 60 87 17 40 98 43 69 48 04 56 62 00 49 49 99 40
81 49 31 73 55 79 14 29 93 71 40 67 53 88 30 03 49 13 36 65 81 49 31 73
52 70 95 23 04 60 11 42 69 24 68 56 01 32 56 71 37 02 36 91 52 70 95 23
22 31 16 71 51 67 63 89 41 92 36 54 22 40 40 28 66 33 13 80 22 31 16 71
24 47 32 60 99 03 45 02 44 75 33 53 78 36 84 20 35 17 12 50 24 47 32 60
32 98 81 28 64 23 67 10 26 38 40 67 55 54 70 66 18 38 64 70 32 98 81 28
67 26 20 68 02 62 12 20 95 63 94 39 63 08 40 91 66 49 94 21 67 26 20 68
24 55 58 05 66 73 99 26 97 17 78 78 96 83 14 88 34 89 63 72 24 55 58 05
21 36 23 09 75 00 76 44 20 45 35 14 00 61 33 97 34 31 33 95 21 36 23 09
78 17 53 28 22 75 31 67 15 94 03 80 04 62 16 14 09 53 56 92 78 17 53 28
16 39 05 42 96 35 31 47 55 58 88 24 00 17 54 24 36 29 85 87 16 39 05 42
86 56 00 45 35 71 89 07 05 46 46 37 44 60 21 58 51 54 17 58 86 56 00 48
19 80 81 68 05 94 47 69 25 73 92 13 86 52 17 77 04 89 55 40 19 80 81 68
04 52 08 83 97 35 99 16 07 97 57 32 16 26 26 79 33 27 98 66 04 52 08 83
88 36 68 87 57 62 20 72 03 46 33 67 46 55 12 32 63 93 33 69 88 36 68 87
04 42 16 73 38 25 39 11 24 94 72 18 08 46 29 32 40 62 76 36 04 42 16 73
20 69 36 41 72 30 23 88 34 62 99 69 82 67 59 85 74 04 36 16 20 69 36 41
20 73 35 29 78 31 90 01 74 31 49 71 48 86 81 16 23 57 05 54 20 73 35 29
01 70 54 71 83 51 54 69 16 92 33 48 61 43 52 01 89 19 67 48 01 70 54 71

What Computers See

Work with Images



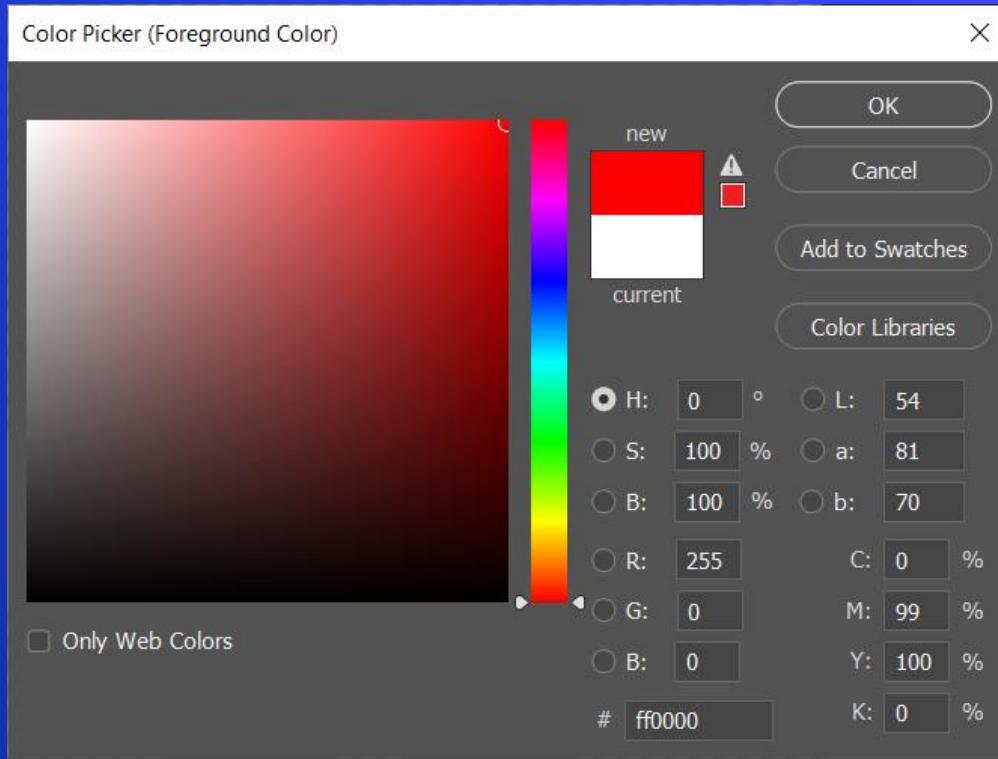
Work with Images

RGB Digital Images

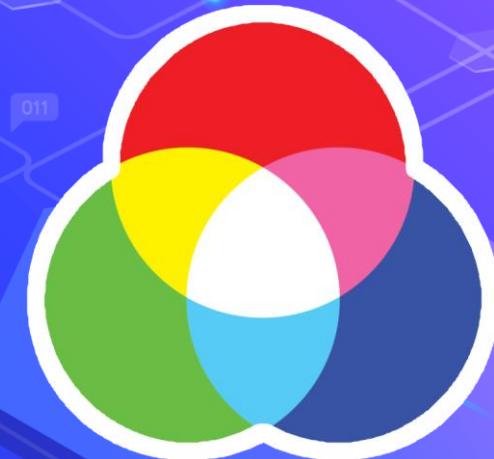
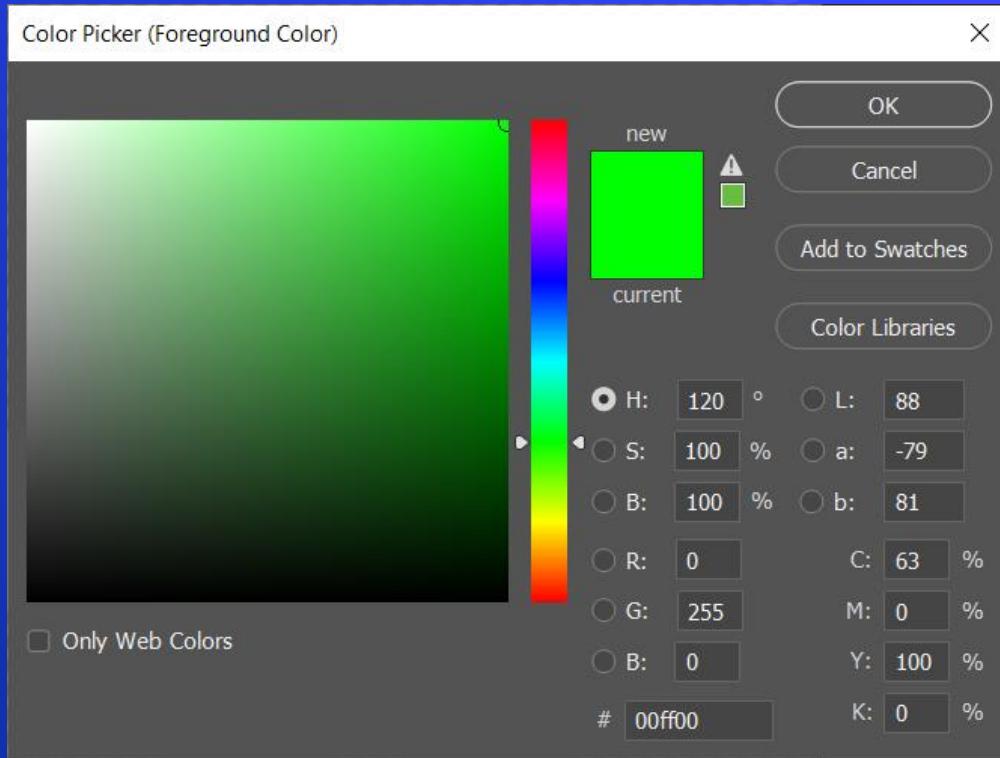
We can generate any color by mixing (red, green, blue) colors with different ratios.



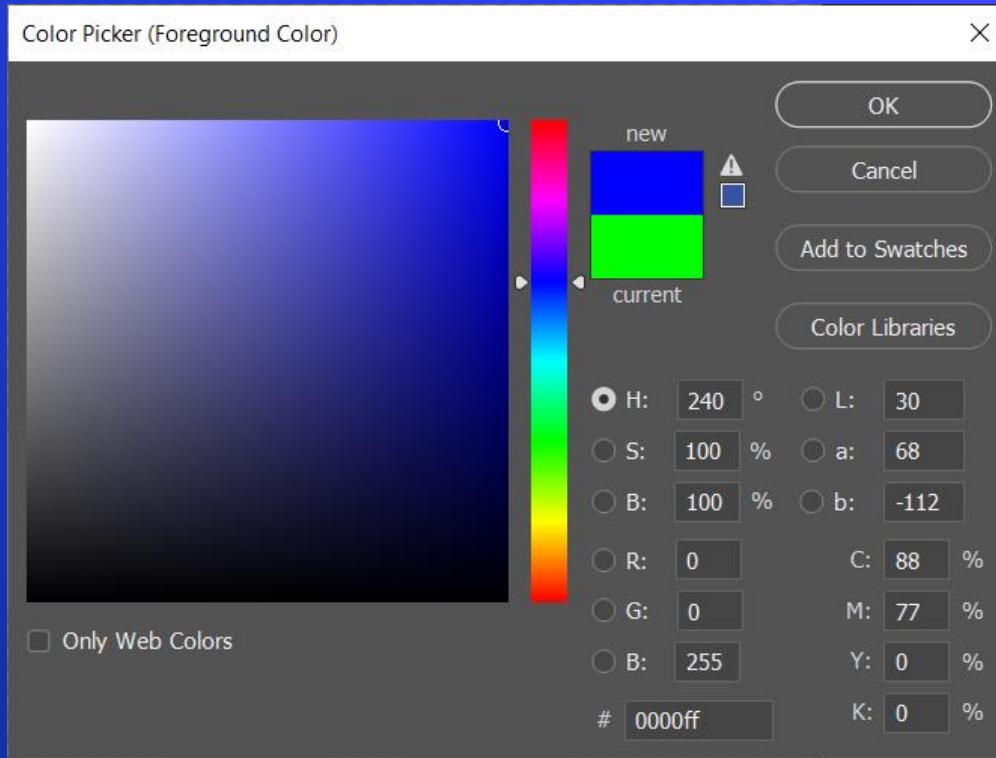
Work with Images (Red)



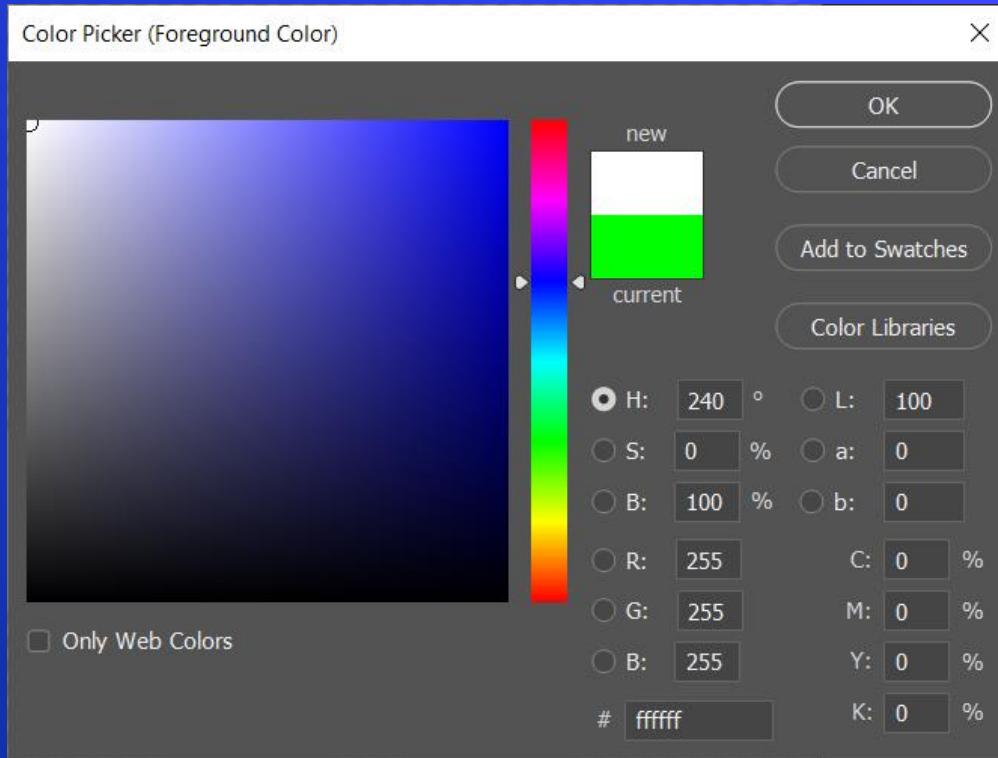
Work with Images (Green)



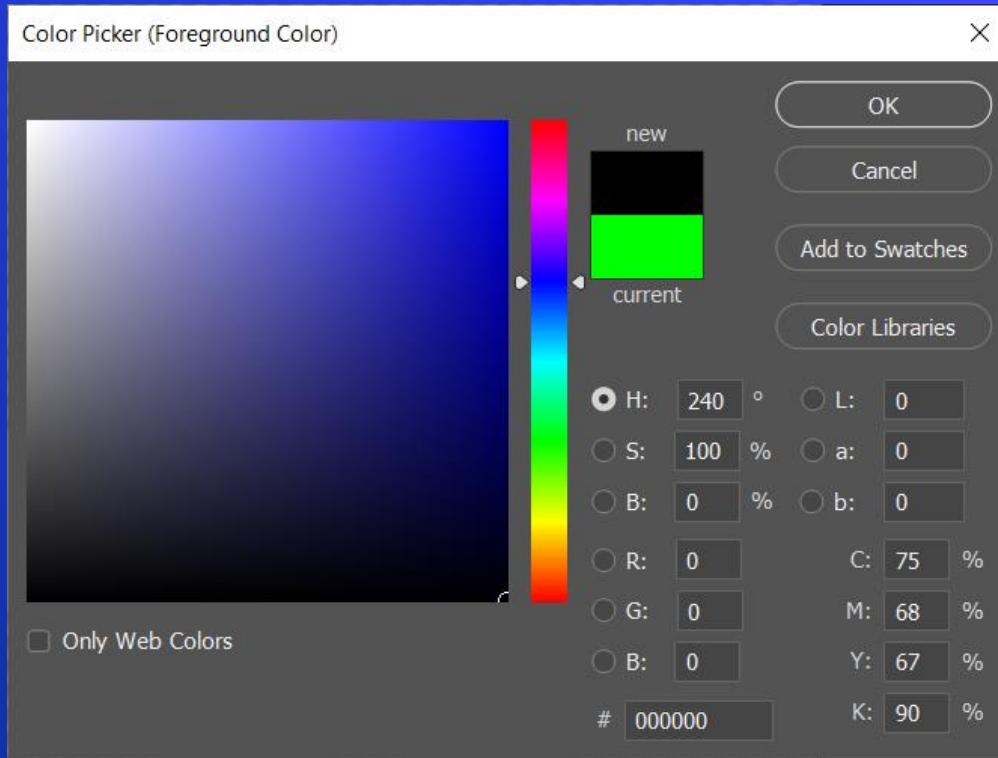
Work with Images (Blue)



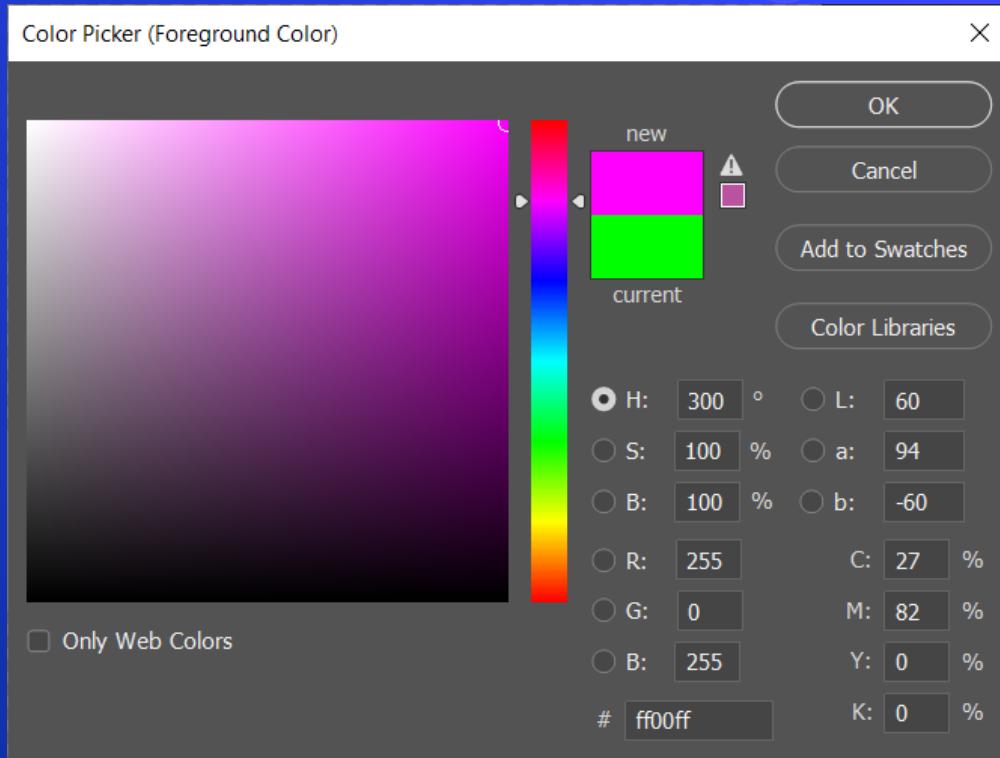
Work with Images (White)



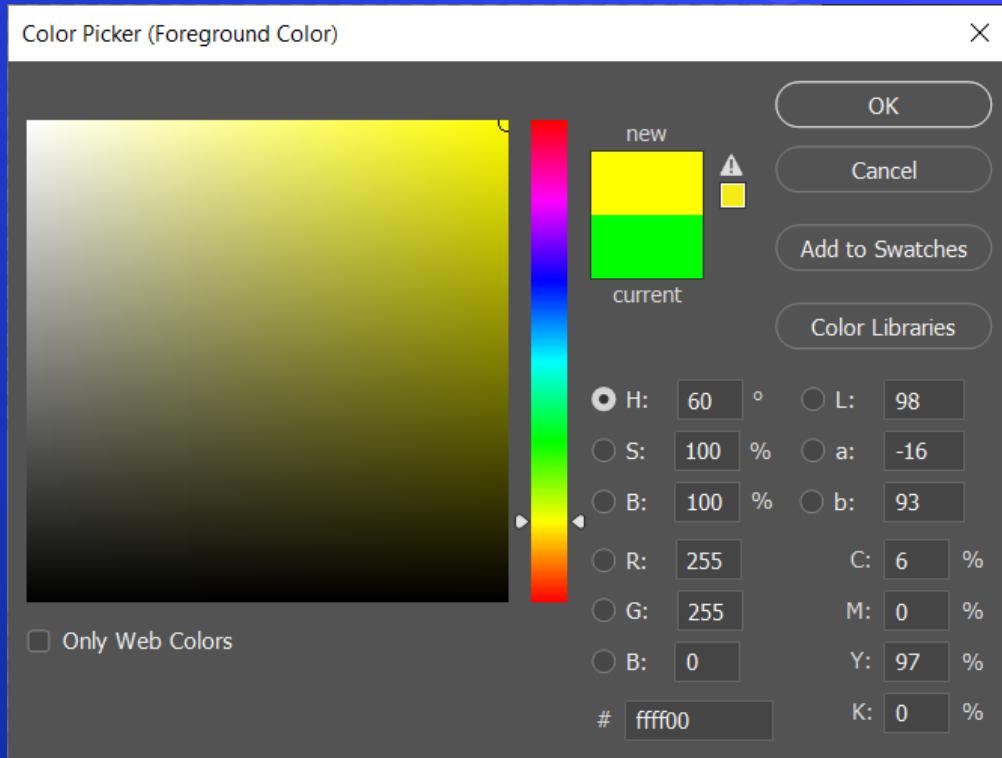
Work with Images (Black)



Work with Images (Pink)



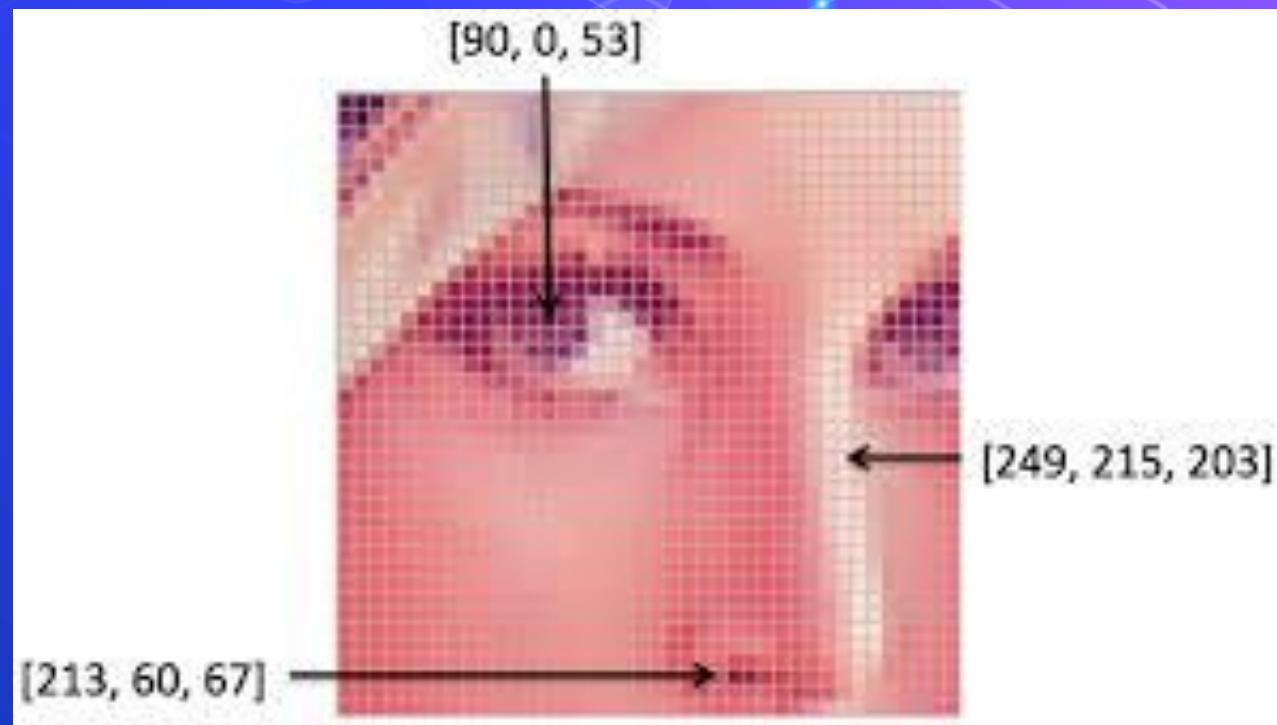
Work with Images (Yellow)



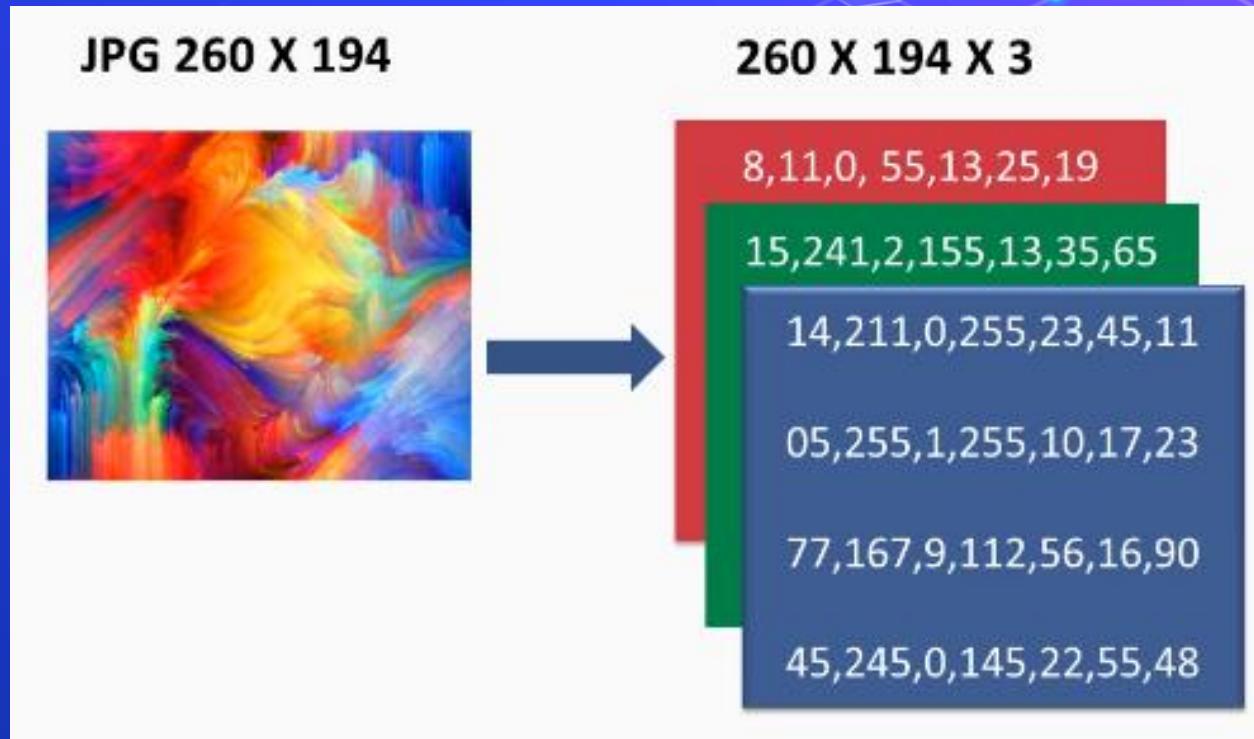
Work with Images

Resolution: 100x100

Width : 100 pixels
Height: 100 pixels



Work with Images

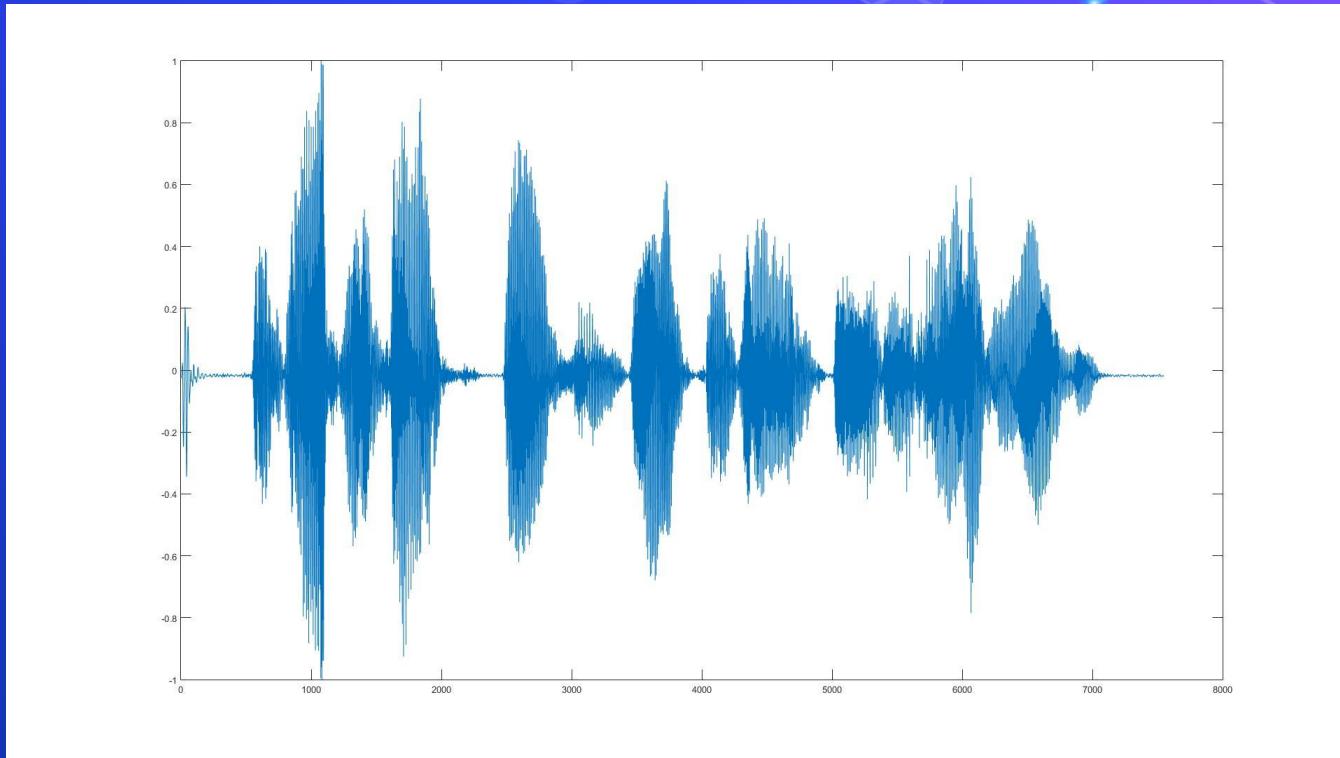


Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - **Work with Audio**
 - Dimension Reduction with PCA

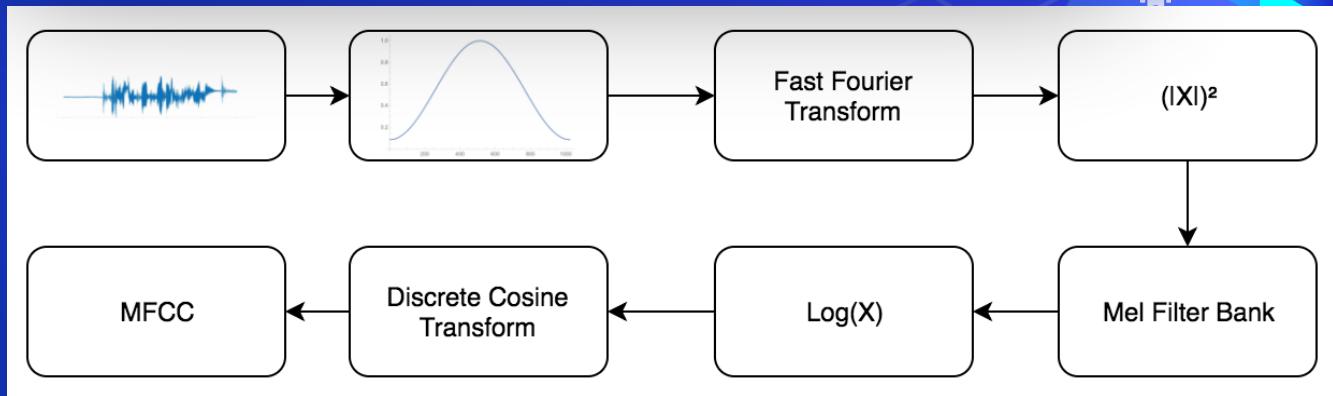
- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Work with Audio



Work with Audio

```
1 import librosa  
2  
3 x, fs = librosa.load('sound.wav')  
4 mfccs = librosa.feature.mfcc(x, sr=fs)
```



https://en.wikipedia.org/wiki/Mel-frequency_cepstrum

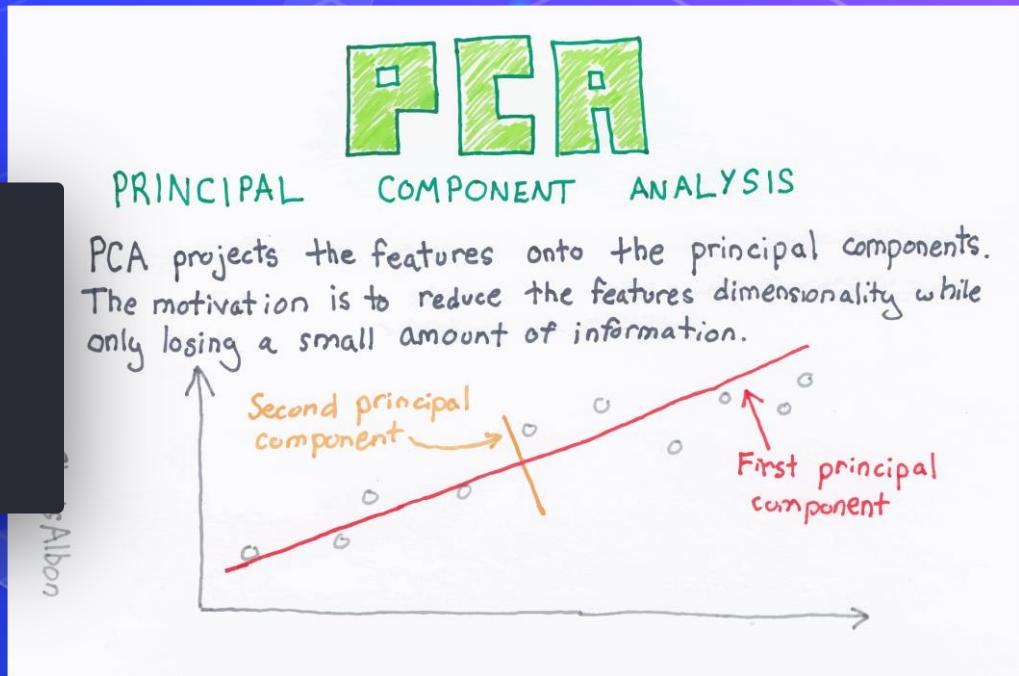
Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon **Feature Engineering and Extraction**
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

- hexagon icon Feature Transformations
 - Data Cleaning or Cleansing
 - Work with Missing data
 - Work with Categorical data
 - Detect and Handle Outliers
 - Split data to Train and Test Sets
 - Deal with Imbalanced classes
 - Feature Scaling
 - How to use Pipelines
- hexagon icon Feature Selection

Dimension Reduction with PCA

```
● ● ●  
1 from sklearn.decomposition import PCA  
2  
3 X = # feature vector  
4  
5 pca = PCA(0.9)  
6 X = pca.fit_transform(X)  
7 pca.explained_variance_ratio_
```



https://en.wikipedia.org/wiki/Principal_component_analysis

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

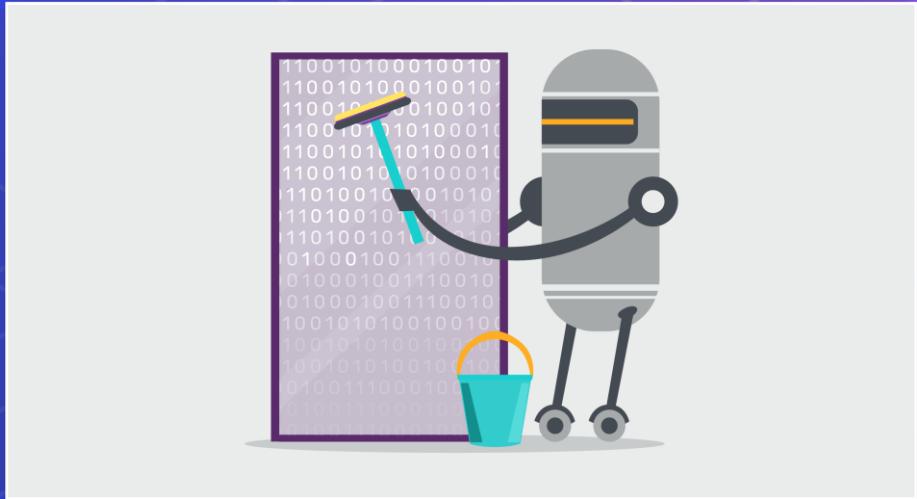
Data Cleaning or Cleansing

Cleaning data is the process of preparing the dataset for analysis. It is very important because the accuracy of machine learning or data mining models are affected because of poor quality of data.

So, data scientists spend a large amount of their time cleaning the dataset and transform them into a format with which they can work with. In fact, data scientists spend 80% of their time cleaning the data.

Example of problems

- Columns can have missing data indicators like xx and ?.
- Height column can have value of 0.
- Weight column can have negative values.
- ...



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

Work with Missing data



```
1 df.isnull().sum()
```

- [Check Missing Data](#)
- [Drop Missing Data](#)
- [Fill Missing data with Pandas](#)
- [Fill Missing data with Sklearn SimpleImputer](#)
- [Fill Missing data with Sklearn KNNImputer](#)



Work with Missing data

```
1 df.dropna(axis = 0)  
2  
3 df.dropna(axis = 1)
```

- Check Missing Data
- [Drop Missing Data](#)
- Fill Missing data with Pandas
- Fill Missing data with Sklearn `SimpleImputer`
- Fill Missing data with Sklearn `KNNImputer`



Work with Missing data

```
1 # fill with mean  
2 df['BuildingArea'].fillna(df['BuildingArea'].mean(), inplace=True)  
3  
4  
5 # fill with median  
6 df['YearBuilt'].fillna(df['YearBuilt'].median(), inplace=True)  
7  
8  
9 # fill with mode (most frequent)  
10 df['Car'].fillna(df['Car'].mode()[0], inplace=True)
```

- Check Missing Data
- Drop Missing Data
- [Fill Missing data with Pandas](#)
- Fill Missing data with Sklearn SimpleImputer
- Fill Missing data with Sklearn KNNImputer



Work with Missing data

```
1 from sklearn.impute import SimpleImputer  
2  
3 # define the Imputer properties  
4 imputer = SimpleImputer(strategy='mean')      # can use 'median' or 'most_frequent'  
5  
6 # impute  
7 df['BuildingArea'] = imputer.fit_transform(df[['BuildingArea']])  
8  
9 # get the filled value  
10 imputer.statistics_
```

- Check Missing Data
- Drop Missing Data
- Fill Missing data with Pandas
- [Fill Missing data with Sklearn SimpleImputer](#)
- Fill Missing data with Sklearn KNNImputer

Work with Missing data

```
1 from sklearn.impute import KNNImputer  
2  
3 imputer = KNNImputer()  
4  
5 df['BuildingArea'] = imputer.fit_transform(df[['BuildingArea']])
```

- Check Missing Data
- Drop Missing Data
- Fill Missing data with Pandas
- Fill Missing data with Sklearn SimpleImputer
- Fill Missing data with Sklearn KNNImputer

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

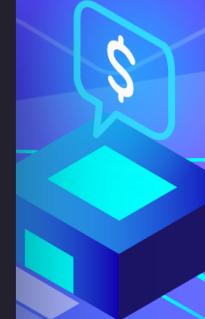
hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

Work with Categorical data

```
1 size_dict = {'XS':1,  
2      'S':2,  
3      'M':3,  
4      'L':4,  
5      'XL':5,  
6      'XXL':6}  
7  
8 # apply using map  
9 df['Size'] = df['Size'].map(size_dict)
```



- Work with Ordinal Features with pandas map method.
- Work with Nominal Features with pandas get_dummies method.

Work with Categorical data

ONE-HOT ENCODING

Feature

Apple
Pear
Apple
Pear
Apple

ONE HOT ENCODING

	Apple	Pear
1	1	0
0	0	1
1	1	0
0	0	1
1	1	0

One-hot encoding allows us to turn nominal categorical data into features with numerical values, while not mathematically implying any ordinal relationship between the classes.

ChrisAlbon



```
1 df = pd.get_dummies(df, columns=['Color', 'Brand'], drop_first=True)
```

- Work with Ordinal Features with pandas `map` method.
- Work with Nominal Features with pandas `get_dummies` method.

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- **Detect and Handle Outliers**
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

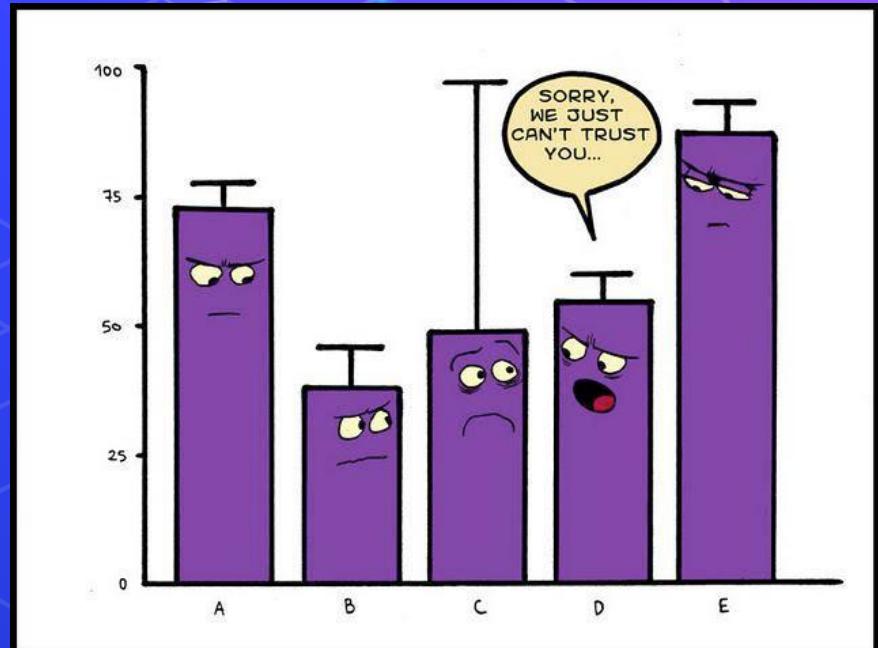
hexagon icon Feature Selection

Detect and Handle Outliers

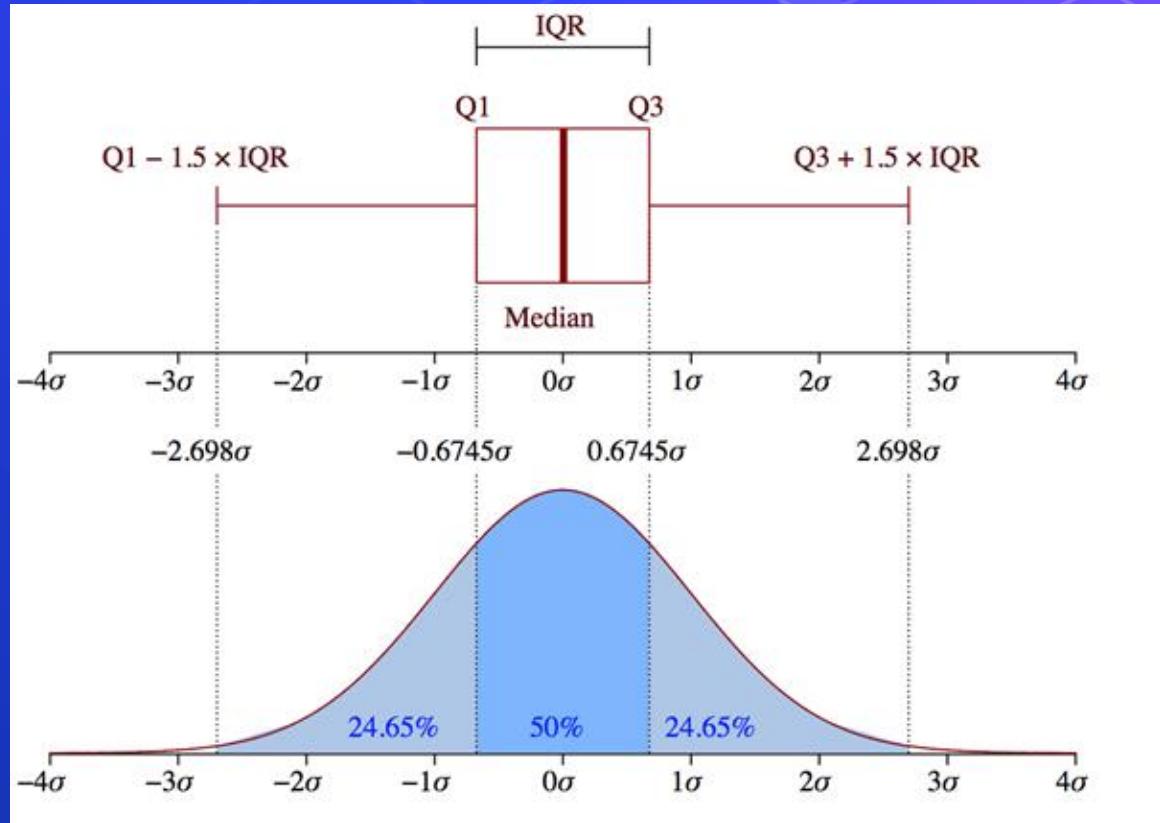
In statistics, an outlier is a data point that differs significantly from other observations.

An outlier may be due to variability in the measurement or it may indicate experimental error, the latter are sometimes excluded from the data set.

An outlier can cause serious problems in statistical analyses.



Detect and Handle Outliers

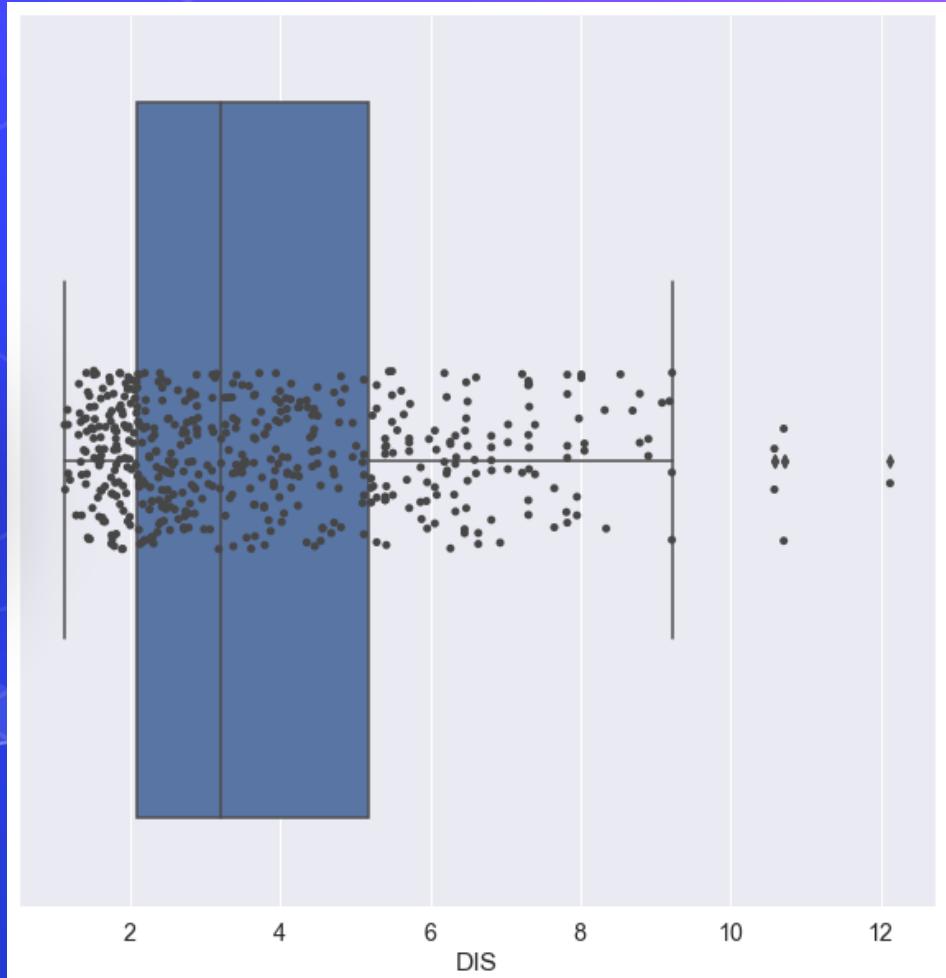


Detect and Handle Outliers



```
1 sns.boxplot(x='DIS', data=df)
2 sns.stripplot(x='DIS', data=df)
```

- [Detect Outliers with Visualization.](#)
- Detect and Handle Outliers with Z-Score.
- Detect and Handle Outliers with IQR.



Detect and Handle Outliers

```
1 # remove outliers that has absolute z-score > 3
2 df = df[(np.abs(stats.zscore(df)) < 3).all(axis=1)]
3
4
5 # replace outliers with median value for each column
6 for col in df.columns:
7     df.loc[np.abs(stats.zscore(df[col])) > 3, col] = df[col].median()
```

- Detect Outliers with Visualization.
- [Detect and Handle Outliers with Z-Score.](#)
- Detect and Handle Outliers with IQR.

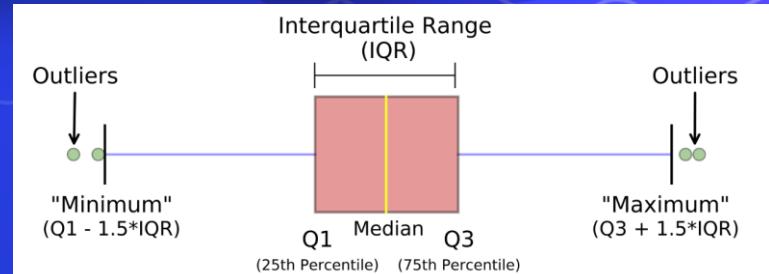
$$z = \frac{(x - \mu)}{\sigma}$$

Data point Mean
 Standard deviation

Detect and Handle Outliers

```
1 # calculate IQR values
2 Q1 = df.quantile(0.25)
3 Q3 = df.quantile(0.75)
4 IQR = Q3 - Q1
5
6
7 # remove outliers < Q1 - 1.5 * IQR or > Q3 + 1.5 * IQR
8 df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
9
10
11 # replace outliers with median value for each column
12 for col in df.columns:
13     rows = (df[col] < (Q1[col] - 1.5 * IQR[col])) | (df[col] > (Q3[col] + 1.5 * IQR[col]))
14     df.loc[rows, col] = df[col].median()
```

- Detect Outliers with Visualization.
- Detect and Handle Outliers with Z-Score.
- Detect and Handle Outliers with IQR.



Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- **Split data to Train and Test Sets**
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

Split data to Train and Test Sets

When you're working on a model and want to train it, you obviously have a dataset. But after training, we have to test the model on some test dataset.

the obvious solution is to split the dataset you have into two sets, one for training and the other for testing; and you do this before you start training your model.

	weight	height	drinks alcohol	healthy
0	112	181	0	0
1	123	165	1	1
2	176	167	1	1
3	145	154	X_train	1
4	198	181	0	0
5	211	202	1	0
6	145	201	1	1
7	181	153	1	1
8	90	142	0	1
9	101	169	X_test	1

Y_train

Y_test

```
1 from sklearn.model_selection import train_test_split  
2  
3 x = df.drop('healthy', axis=1)  
4 y = df['healthy']  
5  
6 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2)
```

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- **Deal with Imbalanced classes**
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

Deal with Imbalanced classes

Imbalanced classes are a common problem in machine learning classification where there are an unbalanced ratio of observations in each class. Class imbalance can be found in many different areas including medical diagnosis, spam filtering, and fraud detection.

Undersampling



Oversampling



Deal with Imbalanced classes

```
1 from sklearn.utils import resample  
2  
3 not_fraud = X[X['Class'] == 0]  
4 fraud = X[X['Class'] == 1]  
5  
6 fraud_upsampled = resample(fraud, replace=True  
                                , n_samples=len(not_fraud))  
7  
8 upsampled = pd.concat([not_fraud, fraud_upsampled])
```

UPSMPLING

A strategy to handle imbalanced classes by repeatedly sample with replacement from the minority class to make it of equal size as the majority class.

ChrisAlbon

- Up-sampling or Over-sampling Minority Class
- Down-sampling or Under-sampling Majority Class
- Generate Synthetic Samples with SMOTE

Deal with Imbalanced classes

```
1 from sklearn.utils import resample  
2  
3 not_fraud = X[X['Class'] == 0]  
4 fraud = X[X['Class'] == 1]  
5  
6 not_fraud_downsampled = resample(not_fraud, replace=False  
7                                     , n_samples = len(fraud))  
8  
9 downsampled = pd.concat([not_fraud_downsampled, fraud])
```

Downsampling

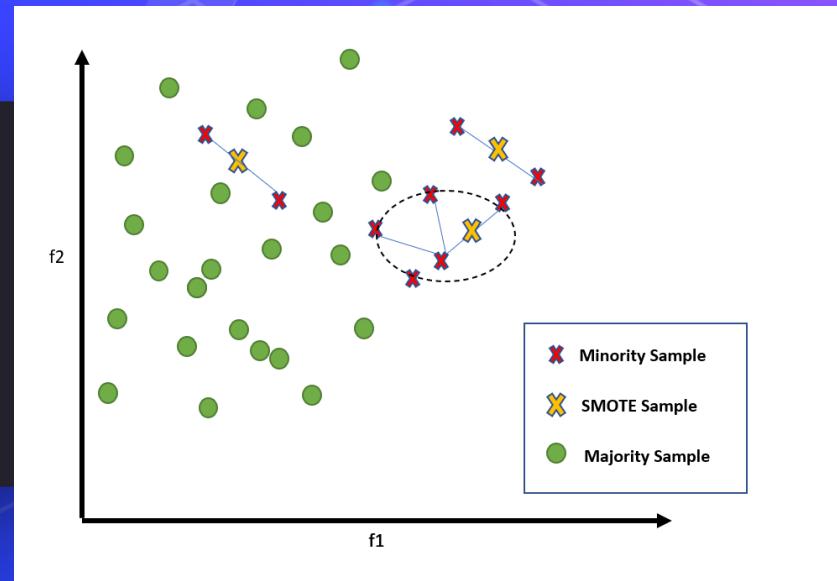
A strategy to handle imbalanced classes by creating a random subset of the majority of equal size to the minority class.

ChrisAlbon

- Up-sampling or Over-sampling Minority Class
- Down-sampling or Under-sampling Majority Class
- Generate Synthetic Samples with SMOTE

Deal with Imbalanced classes

```
● ● ● ● ●  
1 from imblearn.over_sampling import SMOTE  
2  
3 x = df.drop('Class', axis=1)  
4 y = df['Class']  
5  
6 smote = SMOTE()  
7  
8 x_train, y_train = smote.fit_sample(x_train, y_train)
```



- hexagon Up-sampling or Over-sampling Minority Class
- hexagon Down-sampling or Under-sampling Majority Class
- hexagon Generate Synthetic Samples with SMOTE**

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- **Feature Scaling**
- How to use Pipelines

hexagon icon Feature Selection

Feature Scaling

```
● ● ●  
1 from sklearn.preprocessing import MinMaxScaler  
2  
3 scaler = MinMaxScaler()  
4  
5 scaler.fit(x_train)  
6  
7 scaled_x_train = scaler.transform(x_train)  
8 scaled_x_test = scaler.transform(x_test)
```

MIN MAX SCALING

Rescales feature values to between 0 and 1

$$X'_i = \frac{X_i - \min(x)}{\max(x) - \min(x)}$$

Original value Minimum value in feature
Rescaled value Maximum value in feature

ChrisAlbon

- Normalization with Sklearn MinMaxScaler.
- Standardizing data with StandardScaler.
- Rescaling using RobustScaler.

Feature Scaling

•

```
1 from sklearn.preprocessing import StandardScaler  
2  
3 scaler = StandardScaler()  
4  
5 scaler.fit(x_train)  
6  
7 scaled_x_train = scaler.transform(x_train)  
8 scaled_x_test = scaler.transform(x_test)
```

- Normalization with Sklearn MinMaxScaler.
- Standardizing data with StandardScaler.
- Rescaling using RobustScaler.

STANDARDIZATION

$$\tilde{x}_i = \frac{x_i - \bar{x}}{\sigma}$$

Standardized feature value
 \tilde{x}_i ← Value of the i th observation
 \bar{x} ← Mean of the feature vector
 σ ← Standard deviation of the feature vector

Standardization is a common scaling method. \tilde{x}_i represents the number of standard deviations each value is from the mean value. It rescales a feature to have a mean of 0 and unit variance.
Chris Albon

Feature Scaling

```
● ● ●  
1 from sklearn.preprocessing import RobustScaler  
2  
3 scaler = RobustScaler()  
4  
5 scaler.fit(x_train)  
6  
7 scaled_x_train = scaler.transform(x_train)  
8 scaled_x_test = scaler.transform(x_test)
```

ROBUST SCALER

$$\text{scaled value } \tilde{x}_i = \frac{x_i - Q_2}{Q_3 - Q_1}$$

ith observation

scaled value

median or 2nd quartile

1st quartile

3rd quartile

The RobustScaler is very similar to MinMaxScaler.
The difference lies in the parameters used for scaling.
While MinMaxScaler uses minimum and maximum values
for rescaling, RobustScaler uses interquartile(IQR) range
for the same.

- Normalization with Sklearn MinMaxScaler.
- Standardizing data with StandardScaler.
- Rescaling using RobustScaler.

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon Feature Transformations

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon Feature Selection

How to use Pipelines

```
● ● ●  
1 from sklearn.preprocessing import StandardScaler  
2 from sklearn.decomposition import PCA  
3 from sklearn.linear_model import LogisticRegression  
4 from sklearn.pipeline import Pipeline  
5  
6 pipeline_lr = Pipeline([('scalar1',StandardScaler()),  
7                         ('pca1',PCA(n_components=2)),  
8                         ('lr_classifier',LogisticRegression())])  
9  
10 pipeline_lr.fit(X_train, y_train)  
11 pipeline_lr.predict(X_test)
```

Agenda

- hexagon icon What is Data
- hexagon icon Machine Learning
- hexagon icon Data Preprocessing
- hexagon icon Feature Engineering and Extraction
 - Domain knowledge features
 - Date and Time features
 - String operations
 - Web Data
 - Geospatial features
 - Work with Text
 - Work with Images
 - Work with Audio
 - Dimension Reduction with PCA

hexagon icon **Feature Transformations**

- Data Cleaning or Cleansing
- Work with Missing data
- Work with Categorical data
- Detect and Handle Outliers
- Split data to Train and Test Sets
- Deal with Imbalanced classes
- Feature Scaling
- How to use Pipelines

hexagon icon **Feature Selection**

Feature Selection

In machine learning, feature selection is the process of choosing a subset of input features that contribute the most to the output feature for use in model construction.

Feature selection is substantially important if we have datasets with high dimensionality (i.e., large number of features). High-dimensional datasets are not preferred because they have lengthy training time and have high risk of overfitting.

Feature selection helps to mitigate these problems by selecting features that have high importance to the model, such that the data dimensionality can be reduced without much loss of the total information.

Some benefits of feature selection are:

1. Reduce training time
2. Reduce the risk of overfitting
3. Potentially increase model's performance
4. Reduce model's complexity such that interpretation becomes easier

Feature Selection

In filter methods, features are selected independently from any machine algorithms. Filter methods generally use a specific criteria, such as scores in statistical test and variances, to rank the importance of individual features.

- Filter Methods
- Wrapper Methods
- Embedded Methods

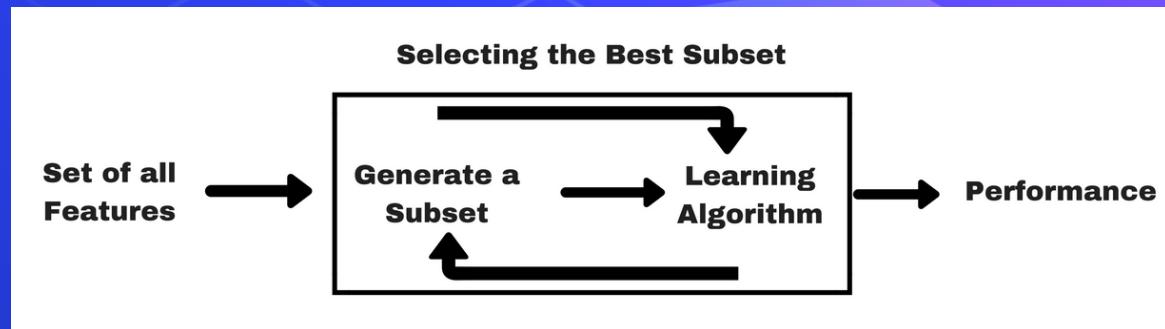


- Pearson's Correlation Coefficient with `f_regression()`
- Mutual Information for regression using `mutual_info_regression()`
- ANOVA for numeric feature using `f_classif()`
- Chi-squared for categorical feature using `chi2()`
- Mutual Information for classification using `mutual_info_classif()`

Feature Selection

Wrapper methods try to find a subset of features that yield the best performance for a model by training, evaluating, and comparing the model with different combinations of features.

- Filter Methods
- Wrapper Methods
- Embedded Methods

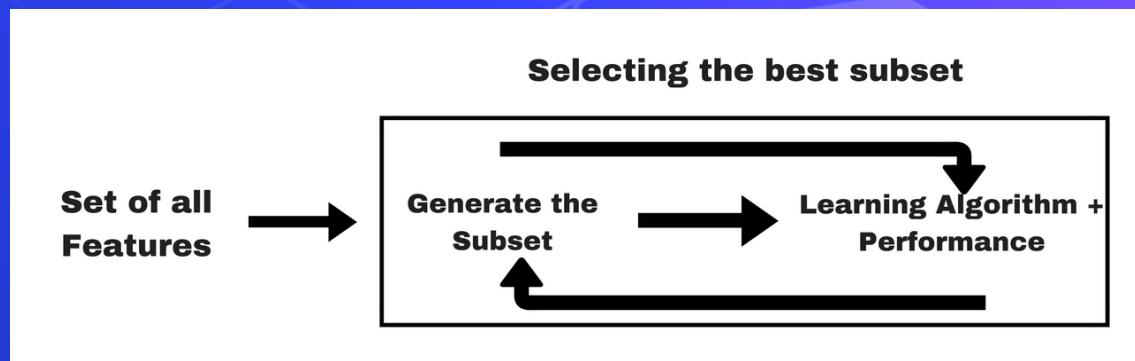


- Recursive Feature Elimination with `RFE()`

Feature Selection

Embedded methods combine the strong points of filter and wrapper methods by taking advantage of machine algorithms that have their own built-in feature selection process. They integrate a feature selection step as a part of the training process

- Filter Methods
- Wrapper Methods
- Embedded Methods



- LASSO & Ridge regression
- Decision trees & Random forests
- Support vector machines

Questions ?!



Thanks!

>_ Live long and prosper

